



**FACULTY
OF MATHEMATICS
AND PHYSICS**
Charles University

MASTER THESIS

Martin Dráb

Thesis title

Name of the department

Supervisor of the master thesis: Petr Daněček

Study programme: study programme

Study branch: Software Systems

Prague 2017

I declare that I carried out this master thesis independently, and only with the cited sources, literature and other professional sources.

I understand that my work relates to the rights and obligations under the Act No. 121/2000 Sb., the Copyright Act, as amended, in particular the fact that the Charles University has the right to conclude a license agreement on the use of this work as a school work pursuant to Section 60 subsection 1 of the Copyright Act.

In date

signature of the author

Title: Thesis title

Author: Martin Dráb

Department: Name of the department

Supervisor: Petr Daněček, department

Abstract: Abstract.

Keywords: key words

Dedication.

Contents

1	Introduction	2
1.1	Basic Terms	2
1.2	The Model	3
1.3	The Real World	4
2	The Algorithm	5
2.1	Input, Output and Preprocessing	5
2.2	Reference transformation	6
2.3	Adding Reads	7
2.3.1	Transforming the Read into K-mers	9
2.3.2	Assigning Sets of Vertices to K-mers	12
2.3.3	The Helper Graph	13
2.4	Graph Structure Optimization	14
2.4.1	Connecting Bubbles	15
2.4.2	Helper Vertices	16
2.5	Variant Detection	17
2.5.1	Simple Bubble	18
2.5.2	Bubble with Inputs	18
2.5.3	Bubble with Outputs	18
2.5.4	Diamond	18
2.6	Varinat Graph and Variant Filtering	19
3	Read Error Correction	20
3.1	De Bruin Graphs	20
3.2	K-mer Frequency Distribution	21
3.3	The Fermi-lite Approach	21
3.3.1	Data Preprocessing	22
3.3.2	Error Correction	23
3.3.3	Unique k-mer Filtering	24
4	Results	26
4.1	Test Data Set	26
4.2	Quality Evaluation	27
4.3	Results	27
4.3.1	Positional	27
	References	29
	List of Figures	30
	List of Tables	31

1. Introduction

Sequencing (reading) a larger genome in one piece still belongs to very expensive, or even hardly possible, tasks. Fortunately, a much cheaper approach exist and is widely used in todays genome sequencing. Its core lies in chopping the genome (or the part of interest) into many short sequences, called reads, and their later assembly into the original sequence. The shorter the reads are, the lower the cost of their reading is.

However, short reads make the task of their assembly more complicated. Unlike the genome chopping phase, the reads assembly stage would execute purely in software if a mathematical model, properly describing the genome sequencing, existed. Since such models were developed, various genome assembly algorithms were implemented.

We start a description of ideal sequencing data and mathematical model behind the assembly. Next, we explain that the real world is not so bright as the model may indicate. FInal section of this chapter is dedicated to description of two main approaches used by todays assembly algorithms, and to definition of the main goal of this work — to implement ou own assembly algorithm.

1.1 Basic Terms

For the rest of this thesis, a genome is viewed as a string of character, each describe a nucleotide, at certain position. Since DNA molecules consist of four types of nucleotides, only four characters, A, C, G or T, appear in the string. Usually, a term base is used as a synonym for nucleotides.

Although we do not know contents of the genome string (since its sequencing is exactly the task for assembly algorithms), we still know its approximate length. Since this thesis focuses on human genome assembly, we expect the genome length around 3 billions of bases (gigabases, Gb). As stated in the beginning of the chapter, the genome to be read is chopped into a large amount of short strings, called reads. In an ideal case, we expect all reads to be of the same length, much lower than the genome string. Usually, read length does not exceed several hundred bases.

For some species, including humans, a so-called *reference* genome (or sequence) is available. Ideally, it is a fully sequenced (assembled) genome of one individual. Since genomes of other individuals of the same species show great similarity, the reference may prove being an useful input for an assembly algorithm.

Some assembly algorithm transform each read into a sequence of *k-mers*, a very short strings of equal length. The read sequence of length k is divided into $l - k + 1$ k-mers k_0, \dots, k_{l-k} of length k . If we denote bases in the read as b_0, \dots, b_{l-1} , the k-mer k_i covers bases from b_i to b_{i+k-1} . Such definition implies that adjacent k-mers overlap by $k - 1$ bases.

An example of transforming e sequence TACTGGCC into k-mers of length 3 is illustrated on Figure 1.1. The sequence has 8 bases in length and 6 k-mers are created from it. The read sequence, as in all other figures in this work, is marked red, contents of individual k-mers is in yellow.




Figure 1.1: Sequence transformation into k-mers

1.2 The Model

The simplest mathematical assembly model, assumes that we are assembling genome string of length g , given a set of n reads of length l ($l \ll g$), each potentially transformed into a sequence of $l - k + 1$ k-mers. The model assumes that bases are, from the genome string, sampled uniformly and randomly and that no errors were produced during the sampling (in other words, all reads contain no misread bases). Since the probability of sampling a base at certain genome position is very low for single sampling event, and the number of sampling events is quite large, the problem of genome base coverage depth (i.e. number of times each base in the genome string appears in reads) follows a Poisson distribution. In other words, the probability that base at certain position is sampled k times is

$$\frac{c^k}{k!} * e^{-c}$$

c is a base coverage depth, also known as sequencing depth, and can be computed simply as a total number of bases within the input read set divided by the length of the genome.

$$c = \frac{n * l}{g}$$

Very similar relations apply for k-mer coverage depth, only the formula for the k-mer coverage depth needs to be changed to $\frac{n * (l - k + 1)}{g - k + 1}$. These details brings answers to at least two important problems: how many reads need to be created to (statistically) cover the whole genome string, and how to determine whether a given read set is free of sequencing errors?

The percentage of genome not covered by any read from the set is equal to $P(k = 0) = e^{-c}$. Multiplying it by the genome size g gives us the number of uncovered bases. So, to cover the whole genome, this number must be lower than 1 which places condition of $c > \ln g$. For example, to cover the whole human genome ($g = 3 * 10^9$), the read coverage depth must be at least 22.

Solution to the second problem is implied by the facts that k-mer coverage of the genome follows a Poisson distribution, and that all the theory above was made with an assumption of no sequencing errors. In an ideal case, the k-mer frequency distribution function of an error-free read set would follow the probability mass function of the Poisson distribution, meaning that frequencies of most of the k-mers are near the k-mer coverage depth. However, when we introduce possibility of sequencing errors, it happens that some k-mers would be sampled less often and some become even unique. The k-mer frequency distribution of such a read set does not follow the Poisson distribution. The aspect of sequencing errors and their means of their correction are described in great detail in Chapter 3.

1.3 The Real World

- give some introduction about importance of genome assembly, as all works of this type usually do,
- define basic terms, such as reference sequence, active region, k-mer, read, base quality, CIGAR string or variant,
- describe current and past approaches for doing genome assembly (AllPath, Euler, Velvet, possibly others. Also HaploCall since it forms a starting point for this thesis).
- formulate the problem this thesis works on,

2. The Algorithm

Our algorithm takes a reference sequence (or a set of them) and a set of reads as inputs, and outputs a VCF file containing all detected variants. The reference sequence must cover the region covered by the read set. Both inputs are read into memory during the initialization phase, there are no memory-saving optimizations employed. In other words, no index files are used for the input data.

The variant calling is done on region basis. The reference sequence is divided into regions of constant length (2000 bases by default), sometimes also referred as *active regions* and with 25 % overlap. Reads are assigned to individual regions according to their mapping position. All regions are called independently and in parallel. To detect the variants, the following steps are performed:

- The reference sequence covering the active region is transformed into a De Bruin-like graph.
- Reads assigned to the active region integrated into the graph.
- The graph structure is optimized and simplified.
- Variants are extracted.
- Variant genotypes and phasing are computed.

2.1 Input, Output and Preprocessing

The reference sequence is expected to be in the FASTA format and starting at position 1. In a typical scenario, the whole chromosome is provided as an input. The FASTA file may contain multiple distinct reference sequences (covering multiple chromosomes). Reference sequences are processed one at a time which means that at most one sequence is present in memory at any moment. Only characters A, C, G and T are considered valid nucleotides. Other characters, such as N, are treated as invalid and reference regions filled with them are not subject to variant calling.

The read set needs to be stored in a text file reflecting the SAM format. Presence of header lines is not required, all information is deduced from the reads. Since the algorithm does not perform any error correction on its own, the input reads must be corrected beforehand. The error correction is supported as a separate command of the tool implementing our algorithm. The correction method was adopted from the **fermi-lite** project and Chapter 3 describes it in great detail.

After the whole SAM file is read into memory and parsed, reads considered useless for the purpose of variant calling are removed from the set. Contents of the **FLAGS** column of the SAM file serves as a main filter, since it is used to detect the following types of undesirable reads:

- **unmapped**, recognized by the bit 2 set to 1,
- **secondary alignments**, detected by the bit 8 set to 1,

- **duplicates**, bit 10 is set to 1,
- **supplementary**, bit 11 is 1.

Also, reads with the mapping position (POS) set to zero and with mapping quality (MAPQ) lower then certain threshold are removed from the set. Read's mapping quality is also used to update individual base qualities. Each base quality q is updated according to the formula

$$b = \min(b, MAPQ)$$

Several other SAM fields play a role in read preprocessing. The CIGAR string is used to detect and remove soft-clipped bases. The QNAME values are used to detect paired end reads, RNAME associates reads into correct input reference sequence.

The SAM file format is described in great detail in *The SAM/BAM format Specification*.

When all the reads are prpeorcessed, their mapping position is used to assign them to individual active regions and the first real phase may begin.

2.2 Reference transformation

The first step of the algorithm lies in transforming a reference sequence covering the selected active region into a De Bruin-like graph. The idea behind this step is very similar to other assembly algorithms based on these graph types.

The reference is divided into k -mers, each overlaps with the adjacent ones by $k - 1$ bases. K -mers representing the same sequence are differentiated by their context number, so each k -mer derived from the reference is unique. Two extra k -mers, denoting the beginning and the end of the active region are added to the set. Then, each k -mer is represented by a single vertex in the graph, and edges are defined by the order of the k -mers within the reference.

Formaly speaking, with the active region of length l represented as a sequence of bases (b_1, \dots, b_l) , k -mers k_0, \dots, k_{l-k+2} are derived from the region as follows:

- $k_0 = ((B, b_1, \dots, b_{k-1}), 0)$
- $k_1 = ((b_1, \dots, b_k), 0)$
- . . .
- $k_i = ((b_i, \dots, b_{i+k-1}), c_i), 2 \leq i \leq l - k + 1$
- . . .
- $k_{l-k+2} = ((b_{l-k+2}, \dots, b_l, E), 0)$

c_i represents the context number of the k -mer k_i . The number is set to zero for k -mers the sequence of which do not repeat within the active region. On the other hand, let's assume that k -mers $k_{i_1}, \dots, k_{i_n}, i_0 < \dots < i_n$ represent the same sequence. Their context numbers are defined as

$$c_{i_j} = j,$$

k_0 and k_{l-k+2} are special k-mers added to the set in order to show the start and end of the active region within the graph. B and E are virtual bases that ensure these k-mers represent unique sequences. The bases must not appear anywhere else within the active region. All k-mers created in this step and all vertices created from them are also called as *reference k-mers* and *reference vertices*. Similarly, k-mers and vertices created during the read integration phase, are sometimes referred as *read k-mers* and *read vertices*.

After their derivation, each k-mer k_i is transformed into a single vertex v_i . Edges follow the order the k-mers occurs within the active region. In other words, the edge set of the graph is

$$E = (v_i, v_{i+1}), 0 \leq i \leq l - k + 1$$

Figure 2.1 displays a graph created by transformation of an active region **ATCTGTATATATG** with k-mer size of 5. The algorithm derives the following k-mers:

$$k_0 = ((B, A, T, C, T), 0) k_1 = ((A, T, C, T, G), 0) k_2 = ((T, C, T, G, T), 0) k_3 = ((C, T, G, T, A), 0) k_4 = ((G, T, A, T, A), 0) k_5 = ((A, T, A, T, A), 0) k_6 = ((T, A, T, A, A), 0) k_7 = ((A, A, T, A, A), 0) k_8 = ((A, A, T, A, A), 0)$$

As can be seen, there are two k-mers representing sequence **TATAT**, namely k_6 and k_8 . Because of their distinct context numbers, they are represented as separate vertices. Introduction of the context numbers removed a loop from the graph. The loop can be observed on Figure 2.2 that shows a standard De Bruin graph constructed from the same active region. K-mers k_6 and k_8 are represented by the same vertex. In order to recover the sequence, it is required to know how many times the loop was actually used during the transformation step.

Although we solved the loop problem, at least for now, by not permitting them to appear, things get more complicated in the next step that covers introducing individual reads into the graph.

2.3 Adding Reads

The basic idea behind this stage is fairly simple and similar to the approach used in the previous one. The read being added is divided into a sequence of k-mers. If a k-mer is not represented by any existing graph vertex, a new vertex is created for it. In other cases, existing vertices are used. Again, vertices representing adjacent k-mers in the read are connected by edges.

Figure 2.3 shows a graph created by transforming a region of **ACCGTGGTAAT** and adding a read **ACCGTAGTAAT** to the resulting graph. K-mer size is set to 5. The read is divided into the following k-mers:

$$k_0 = ((A, C, C, G, T), 0) \tag{2.1}$$

$$k_1 = ((C, C, G, T, A), 0) \tag{2.2}$$

$$k_2 = ((C, G, T, A, G), 0) \tag{2.3}$$

$$k_3 = ((G, T, A, G, T), 0) \tag{2.4}$$

$$k_4 = ((T, A, G, T, A), 0) \tag{2.5}$$

$$k_5 = ((A, G, T, A, A), 0) \tag{2.6}$$

$$k_6 = ((G, T, A, A, T), 0) \tag{2.7}$$




Figure 2.1: Graph resulting from the transformation of ATCTGTATATATG sequence

In the graph, there already are vertices representing k-mers k_0 and k_6 . For others, new vertices are created. Finally, edges are added (if necessary) to show the k-mer order within the read. The figure also suggests how to retrieve the sequence covered by the read — just by following the edges and using the last base of all the k-mers except the first one that is used as a whole.

Figure 2.3 reflects an ideal state, meaning that all places, where the read differs from the reference, are covered by distinct k-mers, and distance between each two of them is greater than k . If these conditions are met, each single n -base long difference (SNP, insertion or deletion) adds atmost $n + k - 1$ new vertices to the graph. However, it may happen that some of the k-mers covering a difference colide by sequence with either k-mers of the reference, or k-mers of other reads covering a totally different place of the active region. To minimize such unfortunate cases, additional graph transformations need to be made after all the reads are added to the graph.

Unfortunately, the basic idea does not work in our case. Introduction of the k-mer context numbers prevented loops in the graph derived purely from the reference. But since multiple k-mers representing the same sequence may exist, it is not always possible to easily determine which of the vertices should be assigned to individual k-mers of the read. For example, if looking at the graph from Figure




Figure 2.2: Transformation of the ATCTGTATATATG sequence to a standard De Bruin graph (with no k-mer context numbers)

2.1, it is not clear which vertex (or vertices) should be assigned to a k-mer with TATAT sequence.

We decided to solve the issue by transforming the basic idea into the following steps:

- divide the read into sequence of k-mers (a so-called *short variant optimization*, described later, may be applied),
- assign a set of vertices to each k-mer, so that all vertices in the set represent k-mers equal to the read k-mer by sequence,
- from each set, select a vertex to represent the k-mer of the read,
- connect the vertices representing the read to respect the order of the k-mers in the read.

2.3.1 Transforming the Read into K-mers

Let's define a read of length n as a sequence (r_1, \dots, r_n) of bases. If the short variant optimization is not applied, the sequence is divided into individual k-mers k_1, \dots, k_{n-k+1} in the same way as for the reference case, except that no extra




Figure 2.3: Basic idea behind adding a read into a De Bruin graph

k-mers to denote read start and end are created. The k-mers look as follows:

$$k_0 = ((r_1, \dots, r_k), 0) \quad (2.8)$$

$$\dots \quad (2.9)$$

$$k_{n-k+1} = ((r_{n_k+1}, \dots, r_n), 0) \quad (2.10)$$

Then, the step described in Subsection 2.3.2 is applied.

As described in ??, in an ideal case, an n -base long difference from the reference produces $n + k - 1$ k-mers different from all reference k-mers. To reduce the probability that some of the k-mers actually colide with either the reference, or another read, the *short variant optimization* may be applied. The optimization reduces the number of k-mers representing a n -base long difference to:

- n for an insertion,
- zero for a deletion,
- 1 for a SNP.

The optimization assumes that when recovering a sequence from the graph, only the last base of each k-mer, with an exception of the starting one, is used. So,

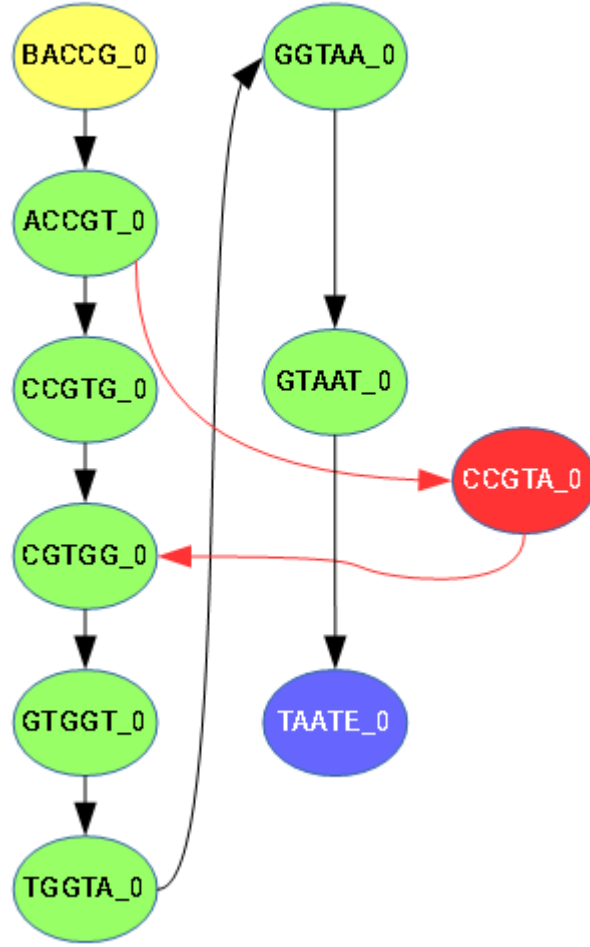


Figure 2.4: Short Variant optimization

only k-mers covering the difference by their last base need to be added; reference k-mers may be used for the rest in case the difference is followed by a reasonable number of bases equal to the reference.

Figure 2.4 shows how the graph is optimized for a read containing SNP. The reference and read sequences are taken from Figure ???. Since the difference has 1 base in length, only one k-mer (CCTGA_0) is used to represent it. The k-mer is followed by reference k-mers. As can be seen, their last base are equal to one of k-mers from Figure ??.

The short variant optimization is applied for k-mer k_i if the following holds:

- There is only one reference vertex with a k-mer equal to k_{i-1} by sequence. Let's assume this is vertex v_{j-1}
- There is at most one vertex for k_i that either is a result of a read addition, or is a reference one but do not immediately follows the vertex v_{j-1} in the reference.

Such conditions are met when the read differs from the reference at base r_{i+k-1} . To determine whether the difference is only a short one, the Smith-Waterman

algorithm is applied. If this is the case, the action depends on the difference type:

- for n -base long deletion, k_i is defined as v_{j+n-1} ,
- for insertion of length n , k_i, \dots, k_{i+n-1} are defined as with no short variant optimization and k_{i+n} is set to v_j ,
- for SNP, k_i is left as such and k_{i+1} is defined as v_{j+1}

2.3.2 Assigning Sets of Vertices to K-mers

The process of assigning vertex sets to individual k-mers derived from the read in the previous step is quite straightforward — a set assigned to certain k-mer x contains exactly the vertices representing k-mers with sequence equal to x . The k-mer context number is not taken into account. If a k-mer is not represented by any vertex of the current graph (thus, the k-mer would receive an empty set), a new vertex is created to represent it. When using a standard De Bruin graph, size of all the sets would contain exactly one vertex and there would be merely anything to speak of. However, introduction of k-mer context numbers caused that also larger sets may appear. That decision, although removing loops from the reference vertices, complicates the task of integrating reads into the graph since the graph may contain multiple path representing a single read (by using different vertices with k-mers equal by sequence).

Previous steps of the algorithm, described above, impose the following conditions on the assigned vertex sets:

- each set contains either read, or reference vertices, not both,
- if a set contains read vertices, its size is always one,
- sets containing reference vertices do not have such restriction,
- each two sets are either distinct (their intersection is an empty set) or equal.

The second and third condition holds because k-mer context numbers are used to differentiate reference k-mers but not the read ones.

In formal terms, a set M_i is assigned to a k-mer k_i where

$$M_i = \{v_{i_j} | v_{i_j} \in V(G), kmer(v_{i_j}) \text{ equals } k_i \text{ by sequence}\}$$

When a set is assigned to each k-mer, it is time to integrate the read into the graph in form of a path, starting on a vertex representing k_1 and ending in one covering k_{n-k+1} . Since M_i sets may contain more than one vertex, it is required to select vertices to form a path best fitting to the read. To derive a good path, we decided to honour the following observations about reads:

- they should follow the reference sequence in a forward direction,
- probability of making big leaps in that direction is low,
- multiple reads cover one place, sharing appropriate parts of their paths.

These observations cannot be enforced too strictly as De Bruin graphs are not very suitable for coping with repeats of length k or more, and similar phenomena. The case of a read containing a copy of reference at least k bases in length might be enough to break the first observation. The second observation permits exceptions by definition. The third one forms a base for most of the genome assembly algorithms.

To solve the problem of selection of the right vertices from the M_i sets, we decided to reduce it to a shortest-path problem on a helper graph the structure of which is defined by the sets and their contents.

2.3.3 The Helper Graph

The helper graph is an oriented layered one. Each layer consists of all vertices contained in one reference M_i sets. The order of the layers respect the order of M_i sets. Sets consisting of read vertices are not part of the helper graph. Only adjacent layers are connected by edges, their orientation reflects the order of the sets. Each subgraph consisting of two adjacent layers is a full bipartite graph. The structure of the helper graph does not take equality of M_i into account. In other words, when $M_i = M_j$ for $i \neq j$, both sets are represented within the helper graph as individual layers, even if they refer to the same vertices of the (main, non-helper) graph.

Formally speaking, let $M_i = \{v_i^1, \dots, v_i^{n_i}\}$ and let i serves as an index to reference sets only. Then the helper graph G_h can be defined as follows:

$$G_h = (V_h, E_h) \quad (2.11)$$

$$V_h = \cup_i M_i \quad (2.12)$$

$$E_h = \{(u, v) | u \in M_i, v \in M_{i+1}\} \quad (2.13)$$

By finding the shortest path leading from a vertex in the first layer to one in the last layer, we perform the process of selection of vertices representing the read in the main graph. The shortest path depends on weights of edges connecting the adjacent layers. In general, the weighting function follows these rules:

- the weight is increased by a *missing edge penalty* if there is a missing edge on the path from u to v in the main graph,
- the weight is increased by a *reference backward penalty* if reference position of u is greater or equal to the reference position of v ,
- the weight is increased by a *reference forward penalty* if reference position of u is far less than reference position of v .

The rules actually indicate why M_i sets covering read vertices are not parts of the helper graph – since their vertices maintain no reference position, only the missing edge penalties would apply and that can be included within missing edge penalties of the reference vertices only.

For an example of a helper graph, let's have a reference sequence **ACTATACTA** and a read **ACTAGACTA**. The left part of Figure ?? shows the main graph just after adding vertices for the reference and the read k -mers with short variant

optimization applied. The resulting helper graph is shown on the right part of the figure.

Six k-mers are derived from the read which means that vertex sets M_0, \dots, M_5 are assigned to them. Since the second k-mer is represented by a read vertex, the M_1 set is not included as a layer of the helper graph. Other sets contains reference vertices, so they form individual layers. Adjacent layers are then connected. Edges with applied penalties (only the reference backward penalty in this case) are depicted red. The black edges show the shortest path.




Figure 2.5: Helper graph creation with short variant optimization

The shortest path select the first vertex from M_0 and the second one from M_5 to represent the read within the main graph (since other sets contain only one vertex, the selection process is trivial there). The resulting path in the graph can be used to correctly recover the sequence covered by the read.

As Figure 2.6 indicates, both graphs look a little bit differently when the short variant optimization is not applied. The main graph contains more read vertices which reduces the number of layers in the helper graph. Although the graphs are different, the sequence covered by the rad is the same and can be correctly recovered again.

2.4 Graph Structure Optimization

When all reads are integrated into the De Bruin-like graph, it is time to optimize its structure in order to get rid of unpopulated paths, usually created by read errors, and resolve some other issues caused mostly by repetitive regions inside either the reference or the reads.




Figure 2.6: Helper graph creation without short variant optimization




Figure 2.7: Benefits of connecting bubbles

2.4.1 Connecting Bubbles

Figure 2.7 demonstrates one class of the structural problems. A subset of reference sequence was transformed into vertices 1, 2, 3, 4, 5 and 6. A set of reads is represented by a "path" 2, 3, 4, 5, R1, R2, 1, 2, 3, R3, 6. The left part of the figure shows how such a graph would look like without any optimizations. It is clear that recovering the correct sequence would not be trivial.

However, since we know that the path leads from R2 to R3 (through 1, 2, 3), we can theoretically replace edges (R2, 1) and (3, R3) by a special edge (R2, R3), as the right part of the figure suggests. An information about the sequence covered by vertices 1 and 2 needs to be recorded within the new edge. Recovering the correct sequence from the right part of the figure does not impose a problem since it is just a simple bubble.

A more general, and a very typical, situation is shown on Figure 2.8. The subgraph contains a subset of the reference (vertices 1, ..., $n+1$) and two bubbles; one ending by $R1$ and connected to 2, another starting at $R2$ and leading from n . If edges I_1 (the input edge) and O_1 (the output edge) share reasonable amount of reads ($|reads(I_1) \cap reads(O_1)| > threshold$), the subgraph may also be interpreted as that the reads contain a sequence of length $n - 1$ that is also present in the reference. In that case, it is wise to connect the vertices $R1$ and $R2$ directly the same way as on Figure ??, bypassing the reference part. The edge maintaining the direct link is marked as C_1 (the connecting edge). Reads shared by the input and the output edges are moved to the connecting one.

If C_1 is created, the read set intersection is also used to decide whether the edges I_1 and O_1 should be removed. The input edge is deleted if it does not share enough reads with the next reference edge (meaning there are no valid sequences

leading through both these edges). Similarly, the output edge is deleted in case it does not share enough reads with the last reference edge (no valid paths goes through the edges).




Figure 2.8: A subgraph required for connection

Figure 2.9 depicts probably the most general case; multiple reads covering different parts of the active region (including the differences from the reference) share the same sequence of n bases. There is k input and l output edges. To determine the association between individual input and output edges, the intersection of covering reads is used again and connecting edges are created if necessary. More precisely, the following rules apply:

- If i^{th} input and j^{th} output edges share reasonable amount of reads ($|reads(I_i) \cap reads(O_j)| > threshold$), a connecting edge $C_{i,j}$ is created and the shared reads are moved to it. The new edge starts in $source(I_i)$ and ends in $dest(O_j)$.
- I_i and O_i are removed in case their read coverage drops below threshold as a result of moving it to the newly created C_i edges.




Figure 2.9: A general form of the subgraph

2.4.2 Helper Vertices

The bubble connection optimization works well when the bubbles being connected contain at least one read vertex. If this is not the case, however, the effect of replacing input and output edges with connection ones leads to a destruction of the sequences recorded within the graph.

As an example, consider a case illustrated by the left part of Figure 2.10. The relevant part of the reference runs from vertex 1 to 7 and the read coverage supports a path of $3 \rightarrow 4 \rightarrow 5 \rightarrow 6 \rightarrow 1 \rightarrow 2 \rightarrow 4 \rightarrow 5 \rightarrow 7$. Applying steps described in this section leads into creating connection edges $6 \rightarrow 4$ and $2 \rightarrow 7$ and removing edges $6 \rightarrow 1$, $2 \rightarrow 4$ and $5 \rightarrow 7$. Such a graph cannot be used to recover the alternate sequence.

Since this problem appears only in case the bubbles are represented only by edges connecting reference vertices (the short variant optimization produces such

cases for deletions), the countermeasure is quite straightforward; it lies in insertion of special vertices that presents themselves as read ones but carry no information about the sequence on which path they exist. We use a conservative approach for their insertion which means they divide the following sorts of edges:

- output degree of their source vertex is greater than one,
- input degree of their destination is greater than one.

The right part of Figure 2.10 indicates where the helper vertices would be inserted. They divide edges $6 \rightarrow 1$, $2 \rightarrow 4$ and $5 \rightarrow 7$. If the bubble connection is applied now, it leads to creation of edges $H1 \rightarrow H2$ and $H2 \rightarrow H3$ and deletion of $H1 \rightarrow 2$, $3 \rightarrow H2$, $H2 \rightarrow 4$ and $5 \rightarrow H3$. The optimization caused the alternate sequence being recoverable ($3 \rightarrow 4 \rightarrow 5 \rightarrow 6 \rightarrow H1 \rightarrow H2 \rightarrow H3 \rightarrow 7$).




Figure 2.10: Use case for helper vertices

2.5 Variant Detection

The graph structure optimization phase ends by removing areas not covered enough by the input read set. Then, alternate sequences covered by the reads are recovered. Since the sequences share most of their parts with the reference, they are extracted in form of variants. Each variant describe one place of the reference where the alternate sequence differ which roughly corresponds to one line of a VCF file.

Variants are extracted by detecting certain subgraphs with preprogrammed interpretations. When such a subgraph is discovered, a variant is deduced from it and integrated back into the graph by replacing its reference edges by one *variant edge*. Read edges unique to the variant are also removed from the subgraph which simplifies its structure. The process of variant detection stops when no suitable subgraphs are found.

Figure 2.11 shows four types of subgraphs used for variant extraction and how the extraction changes them. Requirements on the reference part are always the same; it must consist of a path starting in vertex 1, ending in n and with all inner vertices having only one input and one output edge. Only edges of reference or variant type may be present in the reference part. In all cases, this path is replaced with a variant edge, shown as blue, connecting directly 1 and n . All edges that are removed as a result of the extraction are plotted as discontinuous lines.




Figure 2.11: Variant detection cases

2.5.1 Simple Bubble

Read edges and vertices form a linear path leading from 1 to n . Since the edges are not part of any other variant, they all are removed and the whole subgraph degenerates only to two reference vertices connected by a variant edge. Figure ?? shows this case under the **c)** sign.

2.5.2 Bubble with Inputs

Displayed under the **a)** sign, this sort of subgraphs differ from the simplest one only by allowing input degree of the read vertices to reach above one. However, the restriction on the output degree remains in effect. Only the string of read edges leading from the 1 vertex to the first vertex with more inputs than one are removed. Other read edges may take part also in other variants.

2.5.3 Bubble with Outputs

This subgraph may be viewed as an opposite to the bubble with inputs. All read vertices in the subgraph are allowed to have more outputs than one, but only one input. Only the last string of edges, starting in the last read vertex with output greater than one and ending in the n vertex is removed, since it may participate only in one variant. The case is depicted under the **b)** sign of Figure 2.11.

2.5.4 Diamond

The most complicated case is shown under the **d)** letter and presents sort of a combination of the **a)** and **b)** cases. The sequence of read vertices may contain one with output degree greater than one followed (not necessarily directly) by one with unrestricted input degree. Only the edges covering one part of the supposed diamond are present in one variant only and hence are removed.

2.6 Varinat Graph and Variant Filtering

When variants are extracted from the De Bruin-like graph, they need to be filtered and their genotype and phasing computed. The DB-like graph is not used to help with these tasks. The problem of genotype and phasing is transformed into a sort of graph colouring problem. A so-called *variant graph* is built for this purpose.

The variant graph represent each variant by two vertices; one for its reference and one for its alternate sequence. The final task is to colour each vertex by one of these colors:

- **Blue.** The variant part is used by the first sequence.
- **Red.** The variant part is used by the second sequence.
- **Purple.** Both sequences go through the variant part.

If a vertex is coloured purple, the vertex representing the other part of the variant is not required (no sequence goes through it) and is removed from the graph. The deletion usually happen only to the vertices representing the reference parts, since removing a vertex of the alternate path means that the variant was filtered out.

Before colouring, graph vertices are connected by several types of bidirectional edges that place various conditions on the colour of their sources and destinations.

- **Variant edges** connect vertices representing parts of one variant. Their source and destination must be coloured differently.
- **Read edges** connect variant parts covered by the same subset of reads, such vertices need to be coloured by the same colour, with exception of purple. If one of the vertices is purple, the other may get arbitrary color.
- **Pair edge** put together variant parts that are covered by paired reads. They place the same conditions as read edges.

When the graph is coloured, the genotype and phasing information are known. The variants are written to the resulting VCF file.

3. Read Error Correction

Read data used as an input to many assembly algorithms contain plenty of errors, such as wrongly read bases. To make the data usable for assembly, an error correction step is required. However, it does not remove all the errors and assembly algorithms must cope with that fact, especially when dealing with read ends.

Currently, two different approaches are used to correct read errors, and both are based on transforming individual reads into series of k-mers. One is based on detecting errors as low covered edges (or paths) in a De Bruin graph, the another relies on a k-mer frequency distribution. During development of our algorithm covered in this thesis, we made several attempts to implement an error correction algorithm based on De Bruin graphs. Since we use these graphs also during assembly performing error corrections on them seemed to be a natural choice. Although they definitely improved quality of the input reads, all our attempts did not produce results as good as solutions based on k-mer frequency distribution.

In the end, we decided to adopt the error correction algorithm used by the `fermi-lite` library and based on k-mer frequency distribution. This chapter covers the algorithm in great detail, although it also gives basic information related to usage of De Bruin graphs.

3.1 De Bruin Graphs

This method involves transforming the input set of reads into a De Bruin graph in a way very similar to one used by our assembly algorithm. Although implementation details may differ, the basic idea is the same: each read mapped to certain active region is divided into a sequence of k-mers, each k-mer serves as a vertex and the edges follow the k-mer order within the sequence. Reference sequence, covering the active region, may also be included in the graph.

The most important assumption is that errors produce unique, and thus with low read coverage, connections between graph vertices. Low-covered edges with source vertices that have output degree greater than one are especially interesting. Even a change of a single base in a read sequence may divert a path representing the read through edges with higher read coverage. The locality of the change depends on used k-mer size.

A simple example demonstrating the main idea behind the method is displayed on Figure 3.1. Many reads share a sequence of `TTGCGCTAA`. However, there is also a single read that contains a slightly different one — `TTGCACTAA`. The De Bruin graph shown on the figure uses k-mers 4 bases long. Combination of both sequences produces a standard bubble.

If the bubble was supported by reasonable amount of reads, it would be treated as a SNP. Since only one read supports it, it may be reasonable to consider its divergence from other reads as an error, and to correct it, so the read path would follow more populated edges. When the correction is done, the resulting graph becomes linear, as the right part of Figure 3.1 shows.




Figure 3.1: Simple example of a read error detection by utilizing De Bruijn graphs

3.2 K-mer Frequency Distribution

The method is based on an assumption that k-mer frequency distribution of an error-free read set has certain properties. Especially, frequency of most of the k-mers is from 20 to 40, k-mers with other frequencies are very rare. Figure 3.2 shows the frequency distribution for an error-free read set and for a read set with error rate 1 %.

As can be deduced from the figure, erroneous read sets have less k-mers with frequency between 20 and 40 and contain large amounts of unique or low-frequency ones. The idea behind the correction algorithms based on this method is to transform the low-frequency k-mers in order to move their frequency into the desired interval. Especially k-mers covering bases with low qualities are subjects to changes.

This approach is also used by the **fermi-lite** software and is covered in the next section.

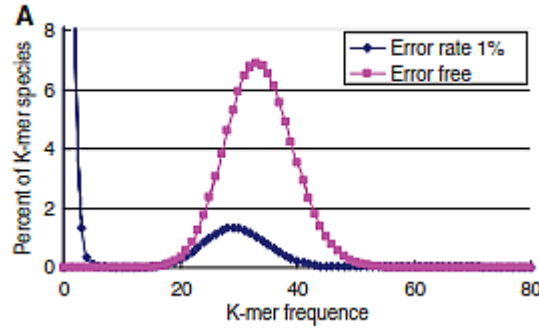


Figure 3.2: K-mer frequency distribution for error-free and error-prone read sets

3.3 The Fermi-lite Approach

Fermi-lite is a standalone C library as well as a command-line tool for assembling Illumina short reads in regions from 100 bp to 10 million bp in size. It is largely a light-weight in-memory version of **fermikit** without generating any intermediate files [from its GitHub]. Results of the assembly are not produced in the VCF format, but rather as a graph. Read error corrections are not the main goal of the project, although this step is definitely required for a successful assembly.

We have successfully extracted the error correction algorithm from the project. The implementation should work well on multiprocessor systems and trades performance over memory consumption. The algorithm proceeds in the following steps:

- **Preprocessing.** The input read sequences are divided into k-mers, k-mer frequencies are calculated.
- **Error correction.** The problem is reduced into a shortest path graph problem and is solved by a kind of Dijkstra algorithm.
- **Unique k-mer filtering.** Unique k-mers introduced during the error correction phase are removed from the read sequences.

3.3.1 Data Preprocessing

The main goal of the preprocessing phase is to compute frequencies for all k-mers found in the input read set. The frequencies are computed by inserting the k-mers into a k-mer table. During this phase, several terms related to k-mers and their occurrences are introduced:

- A k-mer occurrence is defined as *high quality* one if quality of all bases covered by the k-mer is greater than certain threshold. If not all bases satisfy this condition, the occurrence is considered as *low quality*.
- A k-mer is considered *solid* if its frequency is greater than certain threshold.
- A k-mer is considered *unique* if its frequency is zero.

K-mers are implemented as four 64-bit integers,, with structure shown in Figure 3.3. Each base is represented by two bits. That gives limitation for maximum k-mer size to 64. The first 64-bit integer stores least significant bits of the k-mer bases, the second contains their most significant bits. The bases are stored in an opposite order — the first base resides in bits k-1, the second to k-2 and so on. The second two integers store the same content, but with different order — the first base is stored in bits 0, the second in bits 1 etc. Such a k-mer form is redundant; only two 64-bit integers are required to hold all the necessary information. The error correction algorithm actually uses the two-integer representation quite often.

Such a representation allows quick appends or individual base changes. To append a base into the first two integers, they just need to be shifted by one to the left, ORed with the new base, and ANDed with $2^k - 1$ to set the unused bits to zero.

The k-mer table is actually a set of $2^{l_{pre}}$ khash tables. When a k-mer is being inserted or looked up, l_{pre} bits of its data are used to select the table and the rest serves as an input to the hash function. $l_{pre} = 20$ by default. This representation of the k-mer table increases overall memory consumption, but has great impact on its performance in parallel environments.

The table uses 14 bits to track occurrences of each k-mer. Lower 8 bits count low quality occurrences, higher 6 bits are used by high quality ones. The counting stops on values of 255 and 63, no integer overflow happens.




Figure 3.3: K-mer representation used by the `fermi-lite` project

After the insertion stage, a k-mer frequency distribution is computed individually for low quality and high quality occurrences. The most common frequency is named *mode* and is used during the second and third phase of the algorithm.

All phases of the algorithm are partially driven by its parameters. Table 3.1 provides a short description of them.

3.3.2 Error Correction

The error correction is performed separately for each read sequence. The correction problem is transformed into a shortest-path search in a layered oriented graph. Vertices represent individual k-mers, edges connect adjacent ones and their weights reflect the cost of transforming one k-mer to another. The weight function is defined by formula

$$w(k_i) = w_{absent_high} * nhq(k_i) + w_{absent} * nlq(k_i) + w_{ec} * ec(k_i) + w_{ec_high} * bq(k_i) \quad (3.1)$$

$$nhq(k_i) = \begin{cases} 1 & \text{if } k_i \text{ is a high-quality k-mer} \\ 0 & \text{otherwise} \end{cases} \quad (3.2)$$

$$nlq(k_i) = \begin{cases} 1 & \text{if } k_i \text{ is a low-quality k-mer} \\ 0 & \text{otherwise} \end{cases} \quad (3.3)$$

$$ec(k_i) = \begin{cases} 1 & \text{if } k_i \text{ introduces no error correction} \\ 0 & \text{otherwise} \end{cases} \quad (3.4)$$

$$bq(k_i) = \begin{cases} 1 & \text{if the last base of } k_i \text{ has quality above } q \\ 0 & \text{otherwise} \end{cases} \quad (3.5)$$

The search algorithm used is a Dijkstra one and its main loop can be decomposed into the following steps:

- Retrieve the vertex with lowest price, and its k-mer from the heap.
- by separately appending A, C, G, and T to the k-mer, touch the adjacent vertices on the next layer and compute the cost of their connections.
- Insert the newly created vertices into the heap.

Each path from the starting vertex to a vertex in the last layer represents one possible corrected part of the read sequence. Four such paths are computed. The path computation stops if a gap greater than *max_path_diff* is detected in their costs.

3.3.3 Unique k-mer Filtering

The error correction phase may produce unique k-mers which, as Figure 3.2 indicates, are not desirable. The **fermi-lite** library attempts to get rid of such k-mers. Each corrected read sequence is processed separately (and in parallel with others).

At first, the longest k-mer sequence covered by non-unique k-mers is found. Denote its length, in k-mers, as n and the read sequence length as l . Then, the read sequence between the read start and the first base covered by the found k-mer sequence is removed from the read. If the read is covered only by non-unique k-mers, nothing is removed, since the k-mer sequence covers the whole read. However, if the following is true:

$$\frac{n + k - 1}{l} < \textit{min_trim_frac}$$

the read is removed from the read set. This case includes also zero-length k-mer sequences that appear when the read contains unique k-mers only.

Table 3.1: Parameters of the **fermi-lite** algorithm

Parameter	Default value	Description
k	0	K-mer size
q	20	Base quality threshold used to reongize high quality content from low quality one. $P(error) = 10^{\frac{q}{10}}$
min_cov	4	Mimimum frequency for solid k-mers.
win_multi_ec	10	
l_{pre}	20	Number of khash tables in the k-mer table, defined as $2^{l_{pre}}$
min_trim_frac	0.8	Defines how long the sequence of solid k-mers must be in order not to remove the read during unique k-mer filtering.
w_{ec}	1	Participates in a weighing function used in the error correction step.
w_{ec_high}	7	Participates in a weighing function used in the error correction step.
w_{absent}	3	Participates in a weighing function used in the error correction step.
w_{absent_high}	1	Participates in a weighing function used in the error correction step.
max_path_diff	15	Participates in a weighing function used in the error correction step.

4. Results

When an algorithm is being designed, its evaluation against existing solutions belongs to important stages of its development. This chapter describes this stage, informing about a data set used to debug and improve our solution, and the method of comparison with other solutions, such as fermi-kit and GATK. The last part of the chapter covers certain variants that proved to be interesting when examined by our algorithm.

4.1 Test Data Set

The algorithm was tested on the first 40 megabases of chromosome 1 of the human genome. The test set is a high-coverage one and was obtained from the 1000 Genome Project. Except the input reads [1], variants called by fermi-kit and GATK are also available in form of VCF files [3]. The VCF files were used as a measure of algorithm quality. Since our algorithm also requires a reference sequence to work, we took the GRCh37 version [2].

The input read set consists of 12475011 reads with length of 151 bases. Figure ?? shows k-mer frequency distribution of the set with k-mer size of 21 bases. The shape of the graph, when compared to Figure 3.2 suggests that the set contains read errors. Hence, an error correction step was applied.

As Table 4.1 indicates, the error correction process removed and shortened certain amount of reads. About 21 % of the input reads was subject to repairs. Figure 4.1 shows a distribution of the number of repaired bases per read, not including effects of read shortage. It is clear, that about 21 % of all reads received a base correction, and that, in most cases, only several bases were fixed.

Table 4.1: Statistics related to error correction of the test data set

Category	Value	Percentage
Total reads	12475011	-
Removed	64653	0.52 % of all reads.
Shortened	944	0.0076 % of all reads
Total bases	1880123991	-
Bases repaired	5098764	0.27 of all bases

Figure ?? shows the k-mer frequency distribution of the corrected read set. Although quite far from perfect, the graph shape definitely resembles the ideal one better then in case of the raw read set.

As described in Section ??, not all input reads, even from the corrected set, can be processed by our algorithm. Table 4.2 summarizes numbers of reads unacceptable for various reasons. The preprocessing phase removed nearly one fifth of the corrected data set (18.99 %). Most of the reads were removed due to being possible duplicates (87.65 %). Quite a large portion of reads were not accepted because of their low mapping quality (12.97 %). Also, about 3 % of all the reads were shortened in order to remove soft-clipped regions.

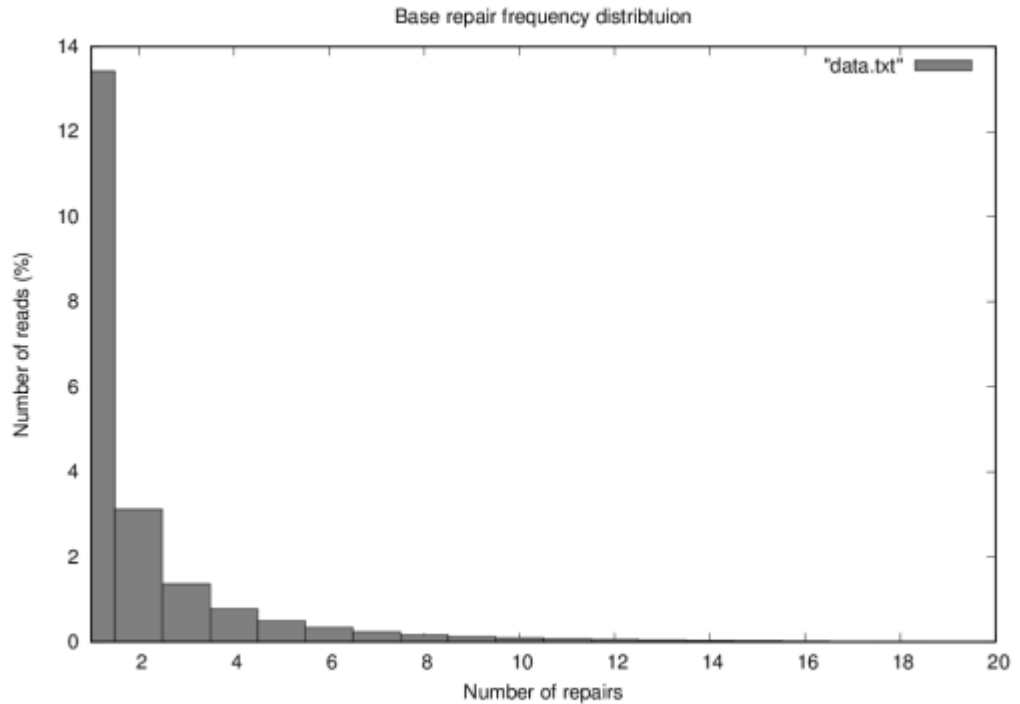


Figure 4.1: Distribution of a number of repaired bases in a single read

4.2 Quality Evaluation

4.3 Results

4.3.1 Positional

- describe algorithms and software to evaluate the results (fermikit, fermikit run at individual regions, GATK, mpileup of samtools, rtgeval for evaluation),
- Show the position-based and genotype results of rtgeval.
- describe interesting variants (variants that are not found by other software, explain some false negatives, demonstrate that graph optimizations actually revealed some variants...).

Table 4.2: Categories of reads present within the corrected test data set

Name	Value	Percentage
Total reads	12410475	-
Bad reads	2357002	18.99 % of all reads
Low MAPQ	305588	12.97 % of bad reads
Unmapped	5020	0.21 % of bad reads
Supplementary	33621	1.43 % of bad reads
Duplicate	2065795	87.65 % of bad reads
Soft-clipped	305209	3.04 % of accempted reads

Table 4.3: Positional comparison of results generated by our algorithm

/	Fermi.kit	Fermi.kit (regions)	mpileup	Our algo
SNP TP	45241 (89,3 %)	47650 (94 %)		48117 (94.96 %)
SNP FN	5432 (10,7 %)	3043 (6 %)		2556 (5.04 %)
SNP FP	385 (0,76 %)	1441 (2,84 %)		803 (1.58 %)
INDEL TP	7853 (71, 6 %)	9802 (89,4 %)		9468 (86.43 %)
INDEL FN	3114 (28,4 %)	1365 (10,6 %)		1499 (13.57 %)
INDEL FP	250 (2,28 %)	835 (7,61 %)		1242 (11.32 %)

References

- [1] ftp://ftp.1000genomes.ebi.ac.uk/vol1/ftp/technical/pilot2_high_cov_GRCh37_bams/data/
- [2] [http : //ftp.1000genomes.ebi.ac.uk/vol1/ftp/technical/reference/human_g1k_v37.fasta](http://ftp.1000genomes.ebi.ac.uk/vol1/ftp/technical/reference/human_g1k_v37.fasta)
- [3] <ftp://ftp.1000genomes.ebi.ac.uk/vol1/ftp/release/20130502/>

List of Figures

1.1	Sequence transformation into k-mers	3
2.1	Graph resulting from the transformation of ATCTGTATATATG sequence	8
2.2	Transformation of the ATCTGTATATATG squence to a standard De Bruin graph (with no k-mer context numbers)	9
2.3	Basic idea behind adding a read into a De Bruin graph	10
2.4	Short Variant optimization	11
2.5	Helper graph creation with short variant optimization	14
2.6	Helper graph creation without short variant optimization	15
2.7	Benefits of connecting bubbles	15
2.8	A subgraph required for connection	16
2.9	A general form of the subgraph	16
2.10	Use case for helper vertices	17
2.11	Variant detection cases	18
3.1	Simple examle of a read error detection by utilizing De Bruing graphs	21
3.2	K-mer frequency distribution for errornous and error-free read sets	21
3.3	K-mer representation used by the fermi-lite project	23
4.1	Distribution of a number of repaired bases in a single read	27

List of Tables

3.1	Parameters of the <code>fermi-lite</code> algorithm	25
4.1	Statistics related to error correction of the test data set	26
4.2	Categories of reads present within the corrected test data set . . .	28
4.3	Positional comparison of results generated by our algorithm . . .	28